

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: data = pd.read_csv("Boston.csv")
data.head()
```

```
Out[3]:   ID      crim     zn  indus  chas    nox      rm    age      dis      rad      tax  ptratio    black
  0   1  0.00632  18.0    2.31      0  0.538  6.575  65.2  4.0900      1  296  15.3  396.90
  1   2  0.02731    0.0    7.07      0  0.469  6.421  78.9  4.9671      2  242  17.8  396.90
  2   4  0.03237    0.0    2.18      0  0.458  6.998  45.8  6.0622      3  222  18.7  394.63
  3   5  0.06905    0.0    2.18      0  0.458  7.147  54.2  6.0622      3  222  18.7  396.90
  4   7  0.08829  12.5    7.87      0  0.524  6.012  66.6  5.5605      5  311  15.2  395.60
```

```
In [4]: data.shape
```

```
Out[4]: (333, 15)
```

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 333 entries, 0 to 332
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   ID          333 non-null    int64  
 1   crim        333 non-null    float64 
 2   zn          333 non-null    float64 
 3   indus       333 non-null    float64 
 4   chas        333 non-null    int64  
 5   nox         333 non-null    float64 
 6   rm          333 non-null    float64 
 7   age         333 non-null    float64 
 8   dis          333 non-null    float64 
 9   rad          333 non-null    int64  
 10  tax          333 non-null    int64  
 11  ptratio      333 non-null    float64 
 12  black        333 non-null    float64 
 13  lstat        333 non-null    float64 
 14  medv        333 non-null    float64 
dtypes: float64(11), int64(4)
memory usage: 39.2 KB
```

```
In [6]: data.isnull().sum()
```

```
Out[6]: ID      0  
        crim     0  
        zn       0  
        indus    0  
        chas     0  
        nox      0  
        rm       0  
        age      0  
        dis      0  
        rad      0  
        tax      0  
        ptratio   0  
        black    0  
        lstat    0  
        medv     0  
        dtype: int64
```

```
In [7]: data.fillna(data.mean(), inplace=True)
```

```
In [8]: X = data[['rm']]  
y = data['medv']
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(  
        X, y, test_size=0.2, random_state=42  
)
```

```
In [10]: single_model = LinearRegression()  
single_model.fit(X_train, y_train)
```

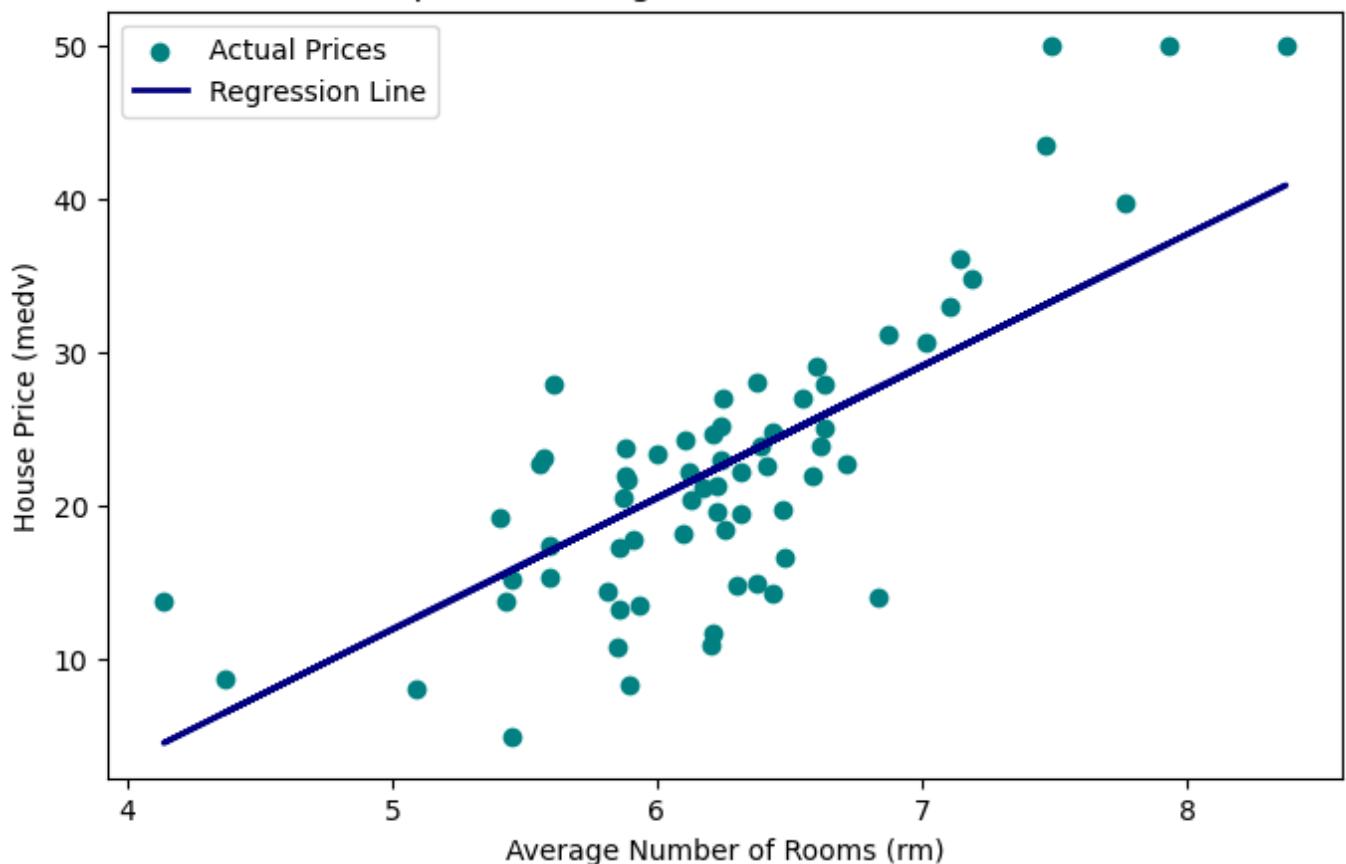
```
Out[10]: ▾ LinearRegression ⓘ ?  
LinearRegression()
```

```
In [11]: y_pred_single = single_model.predict(X_test)  
  
print("Coefficient:", single_model.coef_)  
print("Intercept:", single_model.intercept_)  
print("Mean Squared Error:", mean_squared_error(y_test, y_pred_single))  
print("R2 Score:", r2_score(y_test, y_pred_single))
```

```
Coefficient: [8.58442449]  
Intercept: -30.96185860010203  
Mean Squared Error: 36.361622515889756  
R2 Score: 0.5959747117709422
```

```
In [12]: plt.figure(figsize=(8,5))  
plt.scatter(X_test, y_test, color='Teal', label='Actual Prices')  
plt.plot(X_test, y_pred_single, color='Navy', linewidth=2, label='Regression Line')  
plt.xlabel("Average Number of Rooms (rm)")  
plt.ylabel("House Price (medv)")  
plt.title("Simple Linear Regression: RM vs House Price")  
plt.legend()  
plt.show()
```

## Simple Linear Regression: RM vs House Price



```
In [49]: X_best = data[['tax']]
y = data['medv']

X_train, X_test, y_train, y_test = train_test_split(
    X_best, y, test_size=0.2, random_state=42
)

model_best = LinearRegression()
model_best.fit(X_train, y_train)

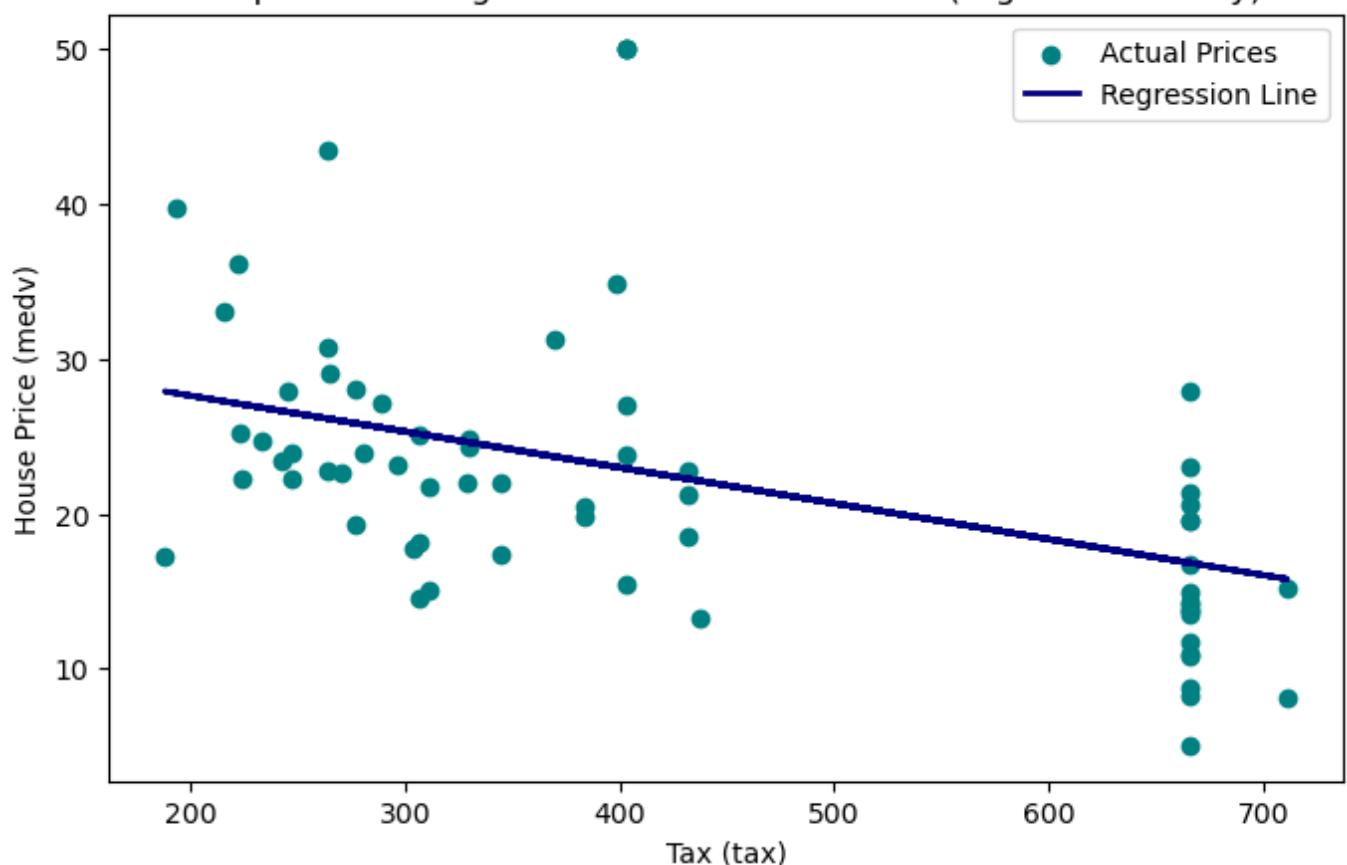
y_pred_best = model_best.predict(X_test)

print("Coefficient:", model_best.coef_)
print("Intercept:", model_best.intercept_)
print("Mean Squared Error:", mean_squared_error(y_test, y_pred_best))
print("R2 Score:", r2_score(y_test, y_pred_best))

plt.figure(figsize=(8,5))
plt.scatter(X_test, y_test, color='Teal', label='Actual Prices')
plt.plot(X_test, y_pred_best, color='Navy', linewidth=2, label='Regression Line')
plt.xlabel("Tax (tax)")
plt.ylabel("House Price (medv)")
plt.title("Simple Linear Regression: RM vs House Price (Highest Accuracy)")
plt.legend()
plt.show()
```

Coefficient: [-0.02316981]  
Intercept: 32.25593991159057  
Mean Squared Error: 67.47367565652733  
R2 Score: 0.25027902033003224

## Simple Linear Regression: RM vs House Price (Highest Accuracy)



```
In [13]: X = data.drop('medv', axis=1)  
y = data['medv']
```

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

```
In [15]: multi_model = LinearRegression()  
multi_model.fit(X_train, y_train)
```

```
Out[15]: ▾ LinearRegression ⚡ ↴  
LinearRegression()
```

```
In [16]: coefficients = pd.DataFrame({  
    'Feature': X.columns,  
    'Coefficient': multi_model.coef_  
})  
  
coefficients
```

Out[16]:

	Feature	Coefficient
0	ID	-0.004902
1	crim	-0.070217
2	zn	0.061337
3	indus	-0.024444
4	chas	4.174129
5	nox	-14.787057
6	rm	3.397953
7	age	-0.014945
8	dis	-1.863079
9	rad	0.421680
10	tax	-0.013171
11	ptratio	-0.734327
12	black	0.008077
13	Istat	-0.637599

In [17]:

```
y_pred_multi = multi_model.predict(X_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred_multi))
print("R2 Score:", r2_score(y_test, y_pred_multi))
```

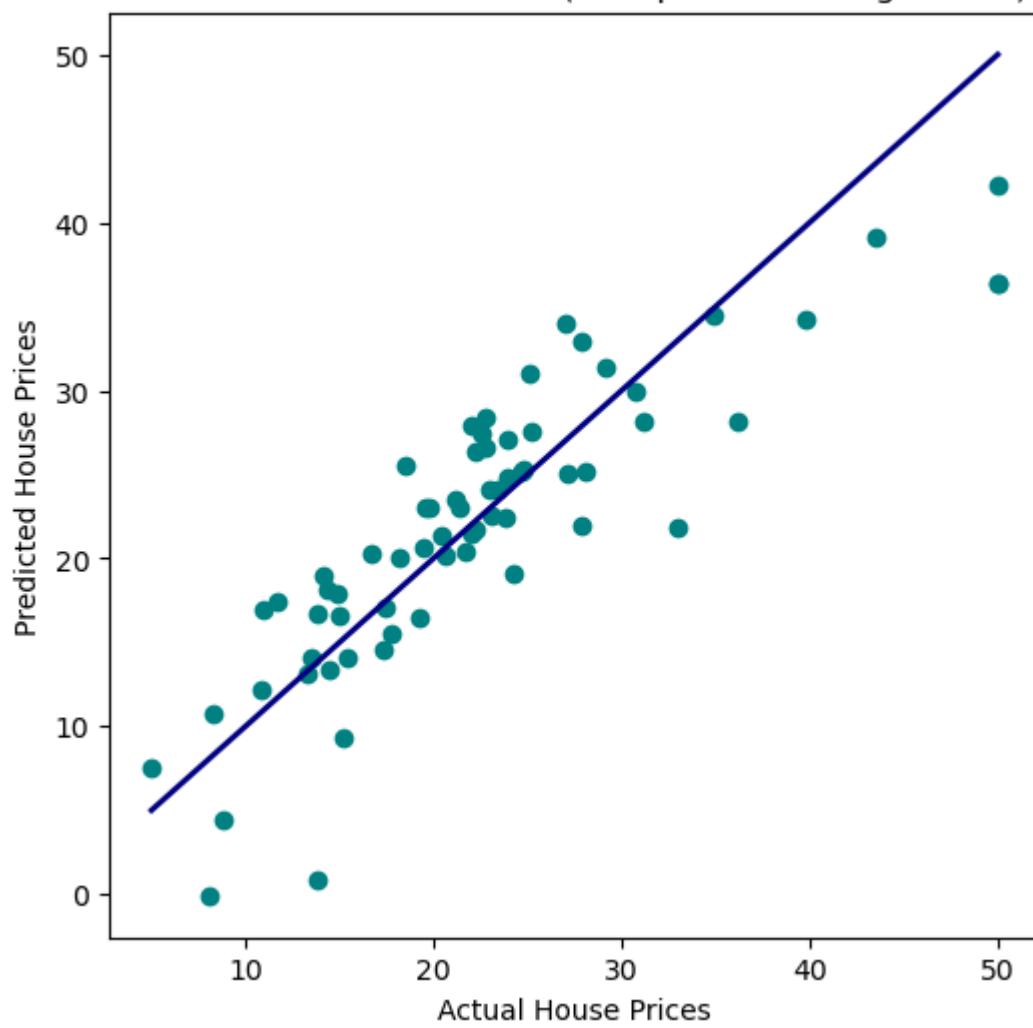
Mean Squared Error: 23.411701826598417

R2 Score: 0.7398653051224349

In [18]:

```
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred_multi, color='Teal')
plt.xlabel("Actual House Prices")
plt.ylabel("Predicted House Prices")
plt.title("Actual vs Predicted Prices (Multiple Linear Regression)")
plt.plot([y_test.min(), y_test.max()],
          [y_test.min(), y_test.max()],
          color='Navy', linewidth=2)
plt.show()
```

### Actual vs Predicted Prices (Multiple Linear Regression)



In [ ]: