

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
```

```
In [4]: data = pd.read_csv("iris.csv")
data.head()
```

```
Out[4]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   sepal length (cm)    150 non-null   float64
 1   sepal width (cm)    150 non-null   float64
 2   petal length (cm)   150 non-null   float64
 3   petal width (cm)   150 non-null   float64
 4   species            150 non-null   object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [6]: data.describe()
```

```
Out[6]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
In [7]: data['species'].value_counts()
```

```
Out[7]: species  
Iris-setosa      50  
Iris-versicolor  50  
Iris-virginica   50  
Name: count, dtype: int64
```

```
In [8]: X = data.drop('species', axis=1)  
y = data['species']
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(  
        X, y, test_size=0.3, random_state=42  
)
```

```
In [10]: nb = GaussianNB()  
nb.fit(X_train, y_train)
```


```
Out[10]: ▾ GaussianNB ⓘ ?  
GaussianNB()
```

```
In [11]: y_pred = nb.predict(X_test)
```

```
In [12]: cm = confusion_matrix(y_test, y_pred)  
cm
```

```
Out[12]: array([[19,  0,  0],  
                 [ 0, 12,  1],  
                 [ 0,  0, 13]])
```

```
In [13]: sns.heatmap(cm, annot=True, cmap='Blues', fmt='d',  
                  xticklabels=nb.classes_,  
                  yticklabels=nb.classes_)  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix")  
plt.show()
```



```
In [14]: for i, label in enumerate(nb.classes_):
    TP = cm[i, i]
    FP = cm[:, i].sum() - TP
    FN = cm[i, :].sum() - TP
    TN = cm.sum() - (TP + FP + FN)

    print(f"\nClass: {label}")
    print("TP:", TP)
    print("FP:", FP)
    print("FN:", FN)
    print("TN:", TN)
```

Class: Iris-setosa

TP: 19
FP: 0
FN: 0
TN: 26

Class: Iris-versicolor

TP: 12
FP: 0
FN: 1
TN: 32

Class: Iris-virginica

TP: 13
FP: 1
FN: 0
TN: 31

```
In [15]: accuracy = accuracy_score(y_test, y_pred)
accuracy
```

Out[15]: 0.9777777777777777

```
In [16]: error_rate = 1 - accuracy
error_rate
```

Out[16]: 0.0222222222222254

```
In [17]: precision = precision_score(y_test, y_pred, average='macro')
precision
```

Out[17]: 0.9761904761904763

```
In [18]: recall = recall_score(y_test, y_pred, average='macro')
recall
```

Out[18]: 0.9743589743589745

```
In [19]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	19
Iris-versicolor	1.00	0.92	0.96	13
Iris-virginica	0.93	1.00	0.96	13
accuracy			0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

```
In [20]: cv_scores = cross_val_score(nb, X, y, cv=5)
cv_scores
```

```
Out[20]: array([0.93333333, 0.96666667, 0.93333333, 0.93333333, 1.])
```

```
In [41]: nb.predict_proba(X_test[:5])
```

```
Out[41]: array([[4.15880005e-088, 9.95527834e-001, 4.47216606e-003],
 [1.00000000e+000, 1.31031235e-013, 2.21772205e-020],
 [9.83170191e-285, 2.70138564e-012, 1.00000000e+000],
 [9.54745274e-092, 9.74861431e-001, 2.51385686e-002],
 [1.08679560e-103, 8.31910700e-001, 1.68089300e-001]])
```

```
In [43]: data = pd.read_csv("spam.csv", encoding='latin-1')
data.head()
```

	Unnamed: 0	label	text	label_num
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n(see...	0
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0

```
In [47]: data.info()
data['label'].value_counts()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5171 entries, 0 to 5170
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0   5171 non-null   int64  
 1   label        5171 non-null   object 
 2   text         5171 non-null   object 
 3   label_num    5171 non-null   int64  
dtypes: int64(2), object(2)
memory usage: 161.7+ KB
```

```
Out[47]: label
ham      3672
spam     1499
Name: count, dtype: int64
```

```
In [49]: X = data['text']
y = data['label_num']
```

```
In [51]: vectorizer = TfidfVectorizer(stop_words='english')
X = vectorizer.fit_transform(X)
```

```
In [53]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)
```

```
In [55]: nb = MultinomialNB()
nb.fit(X_train, y_train)
```

```
Out[55]: ▾ MultinomialNB ⓘ ?
```

```
MultinomialNB()
```


```
In [59]: y_pred = nb.predict(X_test)  
y_pred
```

```
Out[59]: array([0, 1, 0, ..., 0, 0, 0])
```

```
In [61]: cm = confusion_matrix(y_test, y_pred)  
cm
```

```
Out[61]: array([[1121, 0],  
 [130, 301]])
```

```
In [63]: sns.heatmap(cm, annot=True, cmap='Blues', fmt='d',  
                  xticklabels=['Ham', 'Spam'],  
                  yticklabels=['Ham', 'Spam'])  
  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix")  
plt.show()
```



```
In [65]: TP = cm[1,1]  
FP = cm[0,1]  
FN = cm[1,0]  
TN = cm[0,0]
```

```
print("TP:", TP)  
print("FP:", FP)  
print("FN:", FN)  
print("TN:", TN)
```

```
TP: 301  
FP: 0  
FN: 130  
TN: 1121
```

```
In [67]: accuracy = accuracy_score(y_test, y_pred)  
error_rate = 1 - accuracy
```

```
print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
```

```
Accuracy: 0.9162371134020618
Error Rate: 0.08376288659793818
```

```
In [69]: precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print("Precision:", precision)
print("Recall:", recall)
```

```
Precision: 1.0
Recall: 0.6983758700696056
```

```
In [71]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	1.00	0.95	1121
1	1.00	0.70	0.82	431
accuracy			0.92	1552
macro avg	0.95	0.85	0.88	1552
weighted avg	0.92	0.92	0.91	1552

```
In [73]: cv_scores = cross_val_score(nb, X, y, cv=5)

print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

```
Cross-validation scores: [0.91111111 0.93230174 0.92843327 0.92166344 0.91489362]
Mean CV accuracy: 0.9216806361487212
```

```
In [75]: nb.predict_proba(X_test[:5])
```

```
Out[75]: array([[9.99790505e-01, 2.09494966e-04],
 [2.20075521e-01, 7.79924479e-01],
 [9.99998514e-01, 1.48588305e-06],
 [9.99999886e-01, 1.14285146e-07],
 [9.99999879e-01, 1.21435114e-07]])
```

```
In [ ]:
```